



BY Developers FOR Developers

Storage Developer Conference
September 22-23, 2020

ZenFS, Zones and RocksDB

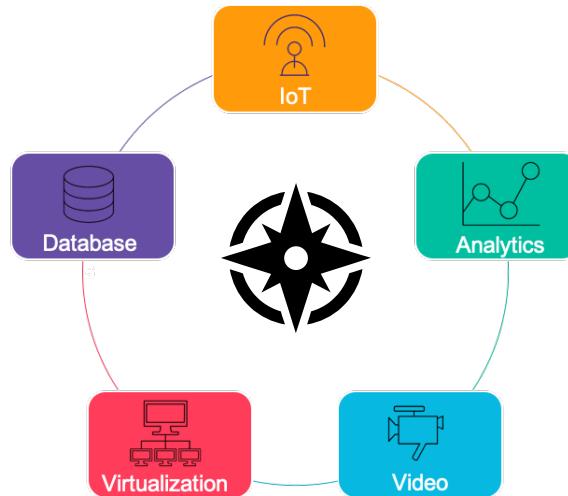
Who likes to take out the garbage anyway?

**Hans Holmberg, Technologist
Western Digital**



Data Growth

- IDC¹ expects that **103 zettabytes** of data will be generated worldwide by 2024
- **How do we scale?**



1) [IDC Worldwide Global DataSphere IoT Device and Data Forecast, 2019–2023](#)

Solid State Drive (SSD)

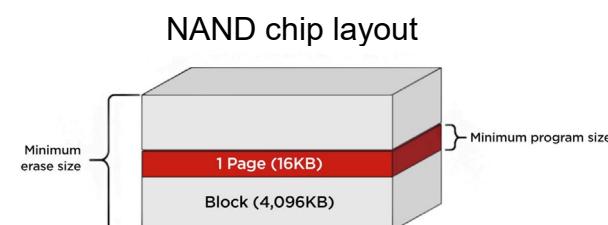
- An SSD bundles 10-100s NAND chips and an SSD controller together
- The SSD controller manages NAND chips characteristics and expose the storage through a storage interface
- A NAND chip is composed of erase blocks, consisting of many pages
 - Within each erase block, you can **only write sequentially**
 - Erase block **must be erased** before new writes
 - Limited number of erases** of an erase block



SSD controller and
10-100x NAND chips



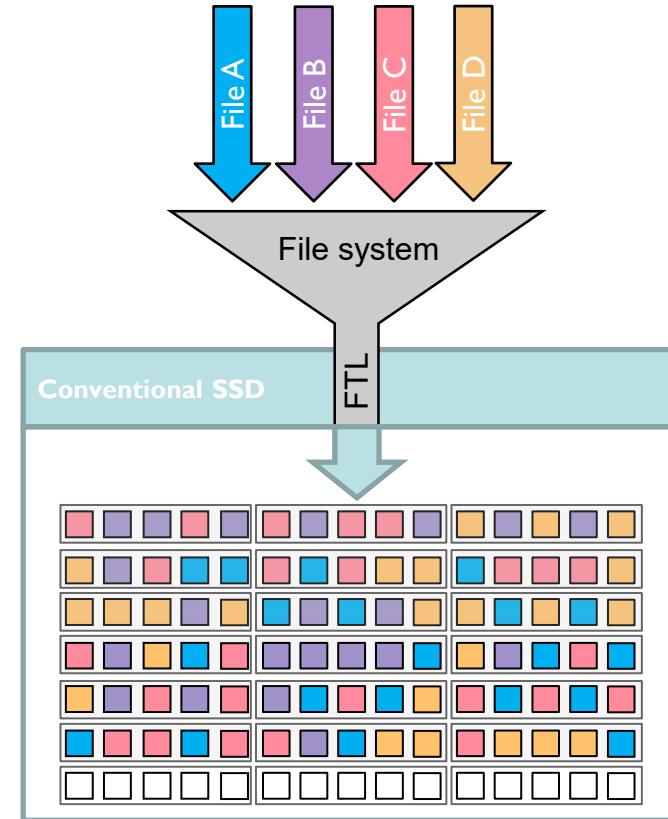
Massive Parallelism



NAND chip layout
Performs a single operation (Read/Program/Erase) at a time

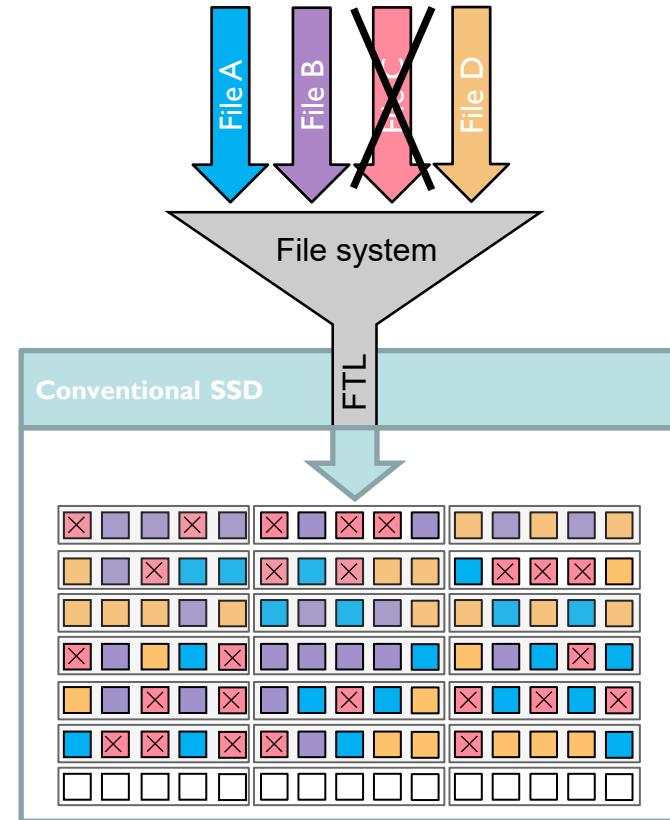
Conventional SSDs

- A Flash Translation Layer (FTL) maps logical blocks to physical addresses
- Files can't be separated by the drive
- Different types of data gets co-located in the same erase units



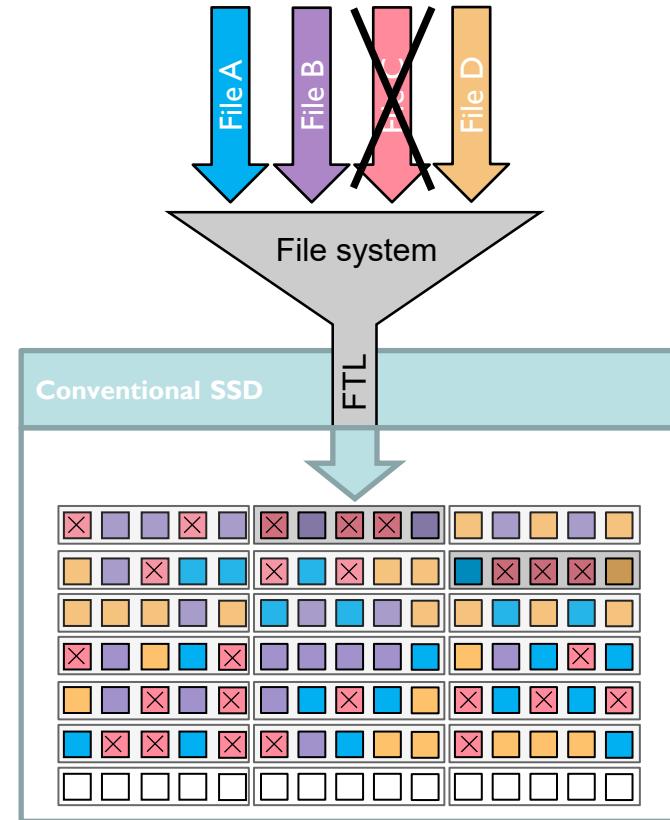
Space reclaim

- When a file gets deleted the space occupied can't be reused without erasing an entire erase unit
- Garbage collection is needed to evacuate still-valid data from the erase unit



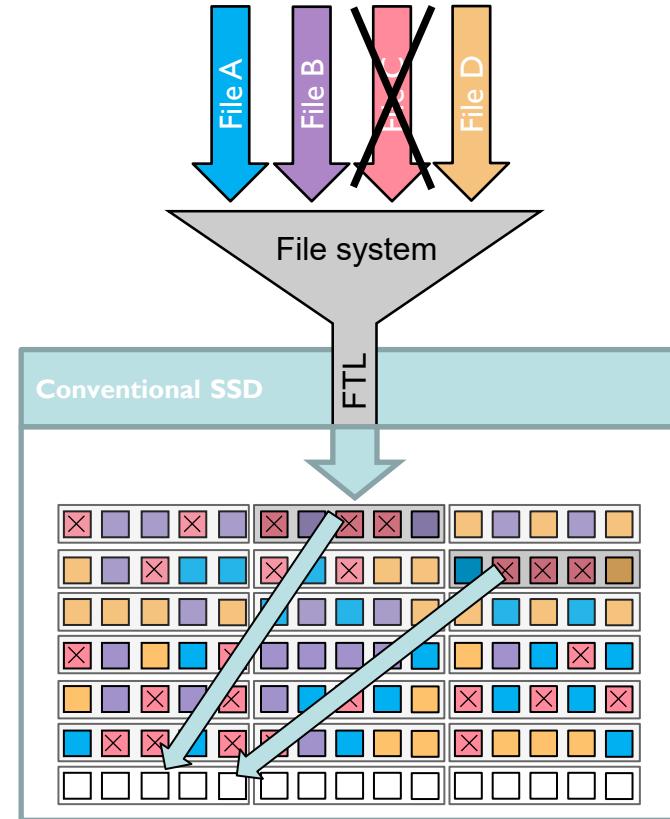
Space reclaim

- The controller picks erase units to be evacuated



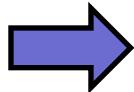
Space reclaim

- The controller picks erase units to be evacuated
- Still-valid data is evacuated to another erase unit



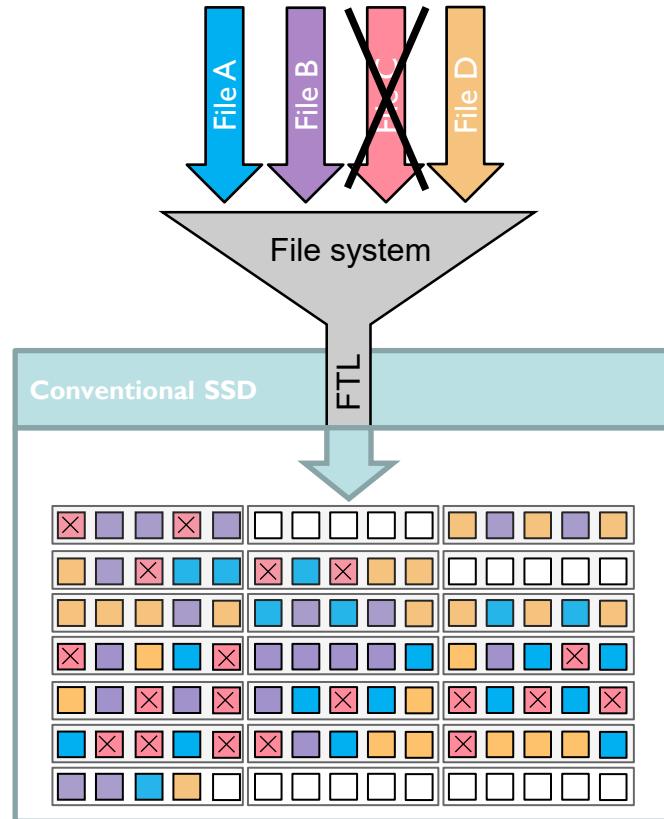
Space reclaim

- Once an erase unit is evacuated it can be erased and reused
- Garbage collection causes rewrites of user data



Write Amplification

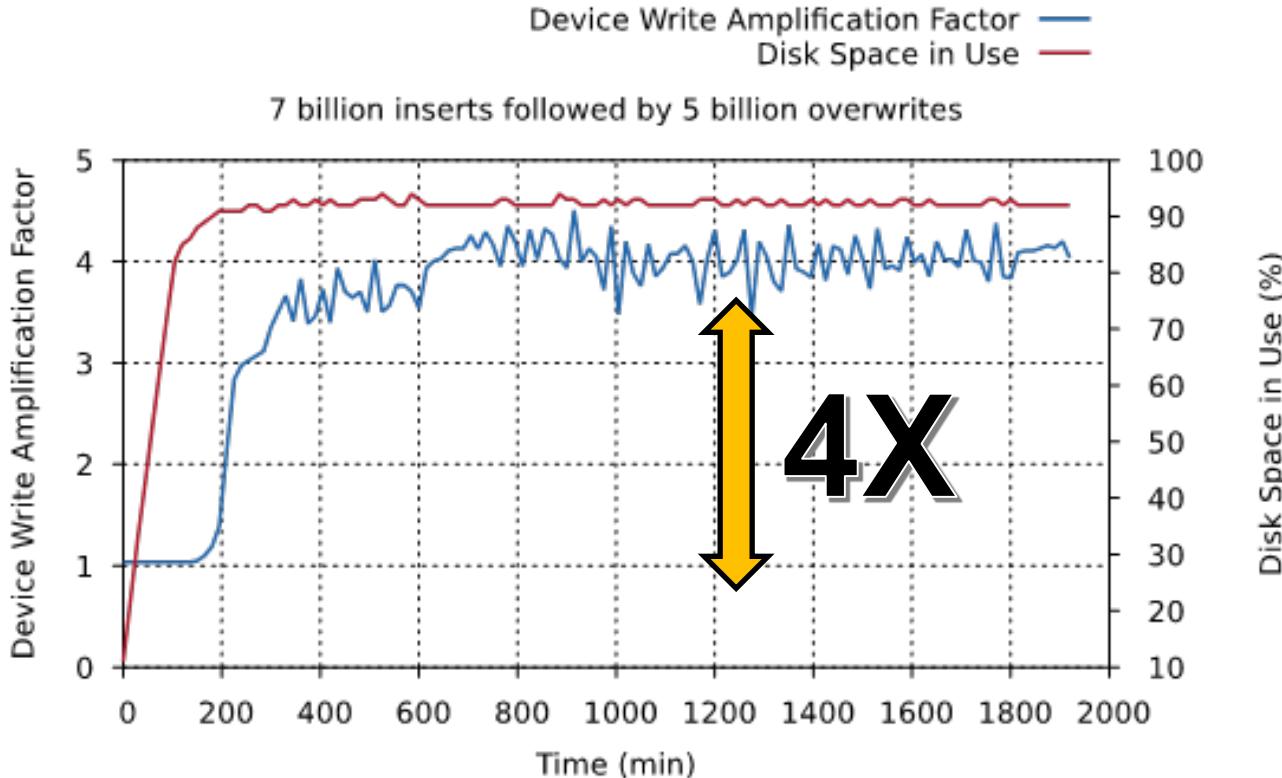
$$WA = \frac{\text{drive blocks written}}{\text{host blocks written}}$$



Write Amplification

- More drive writes per user write
 - Shorter life span of the drive
 - Decreased write performance
 - Worse read tail latencies

RocksDB WA on a conventional SSD



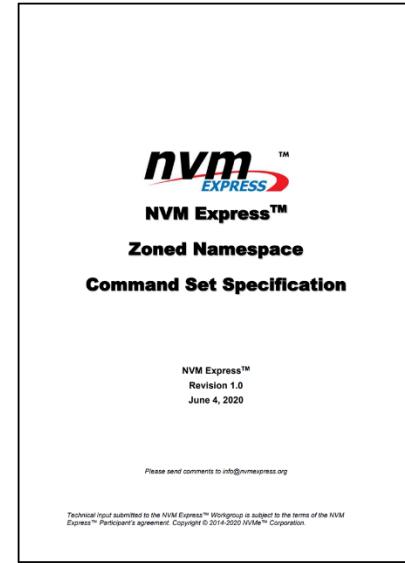


**Who likes to take out the
garbage anyway?**

Can't we do better?

NVMe™ Zoned Namespace Command Set

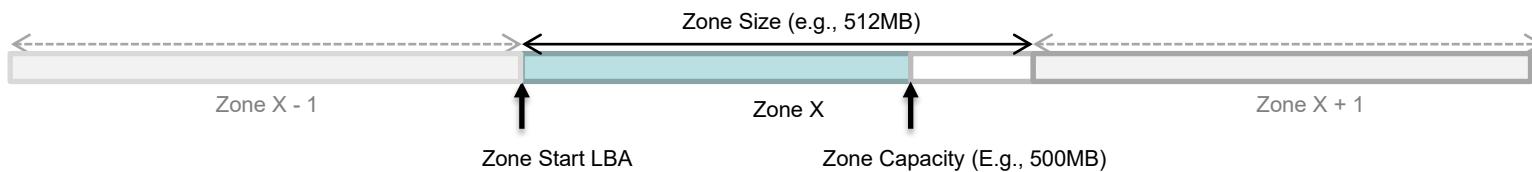
- Introduces the Zoned Storage Model for NVMe™
- Introduces a new namespace type (Zoned Namespaces)
 - Exposes a set of zones of fixed size to be written sequentially and reset for new writes (matches the NAND media characteristics)
 - Implements the Zoned Namespaces Command Set
 - The command set inherits the NVM Command Set
 - i.e., Read/Write/Flush commands are available.
- Optimized for SSDs
 - Adds a more efficient interface for flash media
 - Improves system-level performance



<https://nvmexpress.org/developers/nvme-specification/>
(Available in the 1.4 TP package)

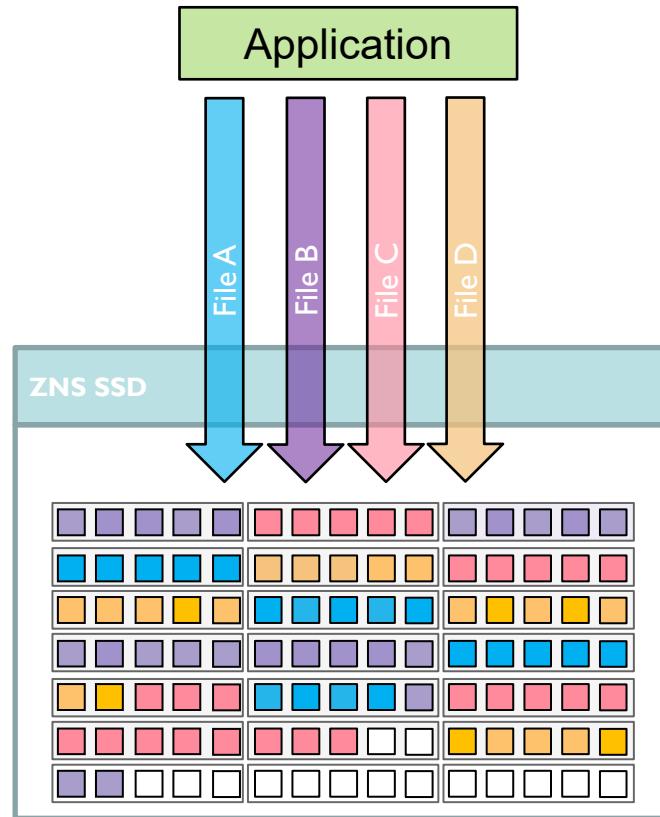
Zoned Storage Model Overview

- Zone States
 - Empty / Open / Closed / Read Only / Offline
- Zone Management
 - Open Zone, Close Zone, Finish Zone, and Reset Zone
- Zone Size & Zone Capacity
 - Zone Size is fixed
 - **Zone Capacity** is the writeable area within a zone
- **Active** and Open Resources associated to a zone
 - Limits the maximum active and open zones



ZNS enables smart data placement

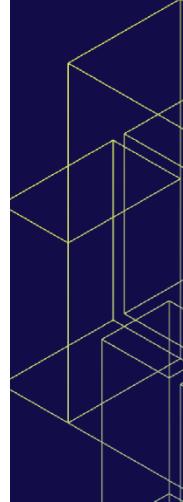
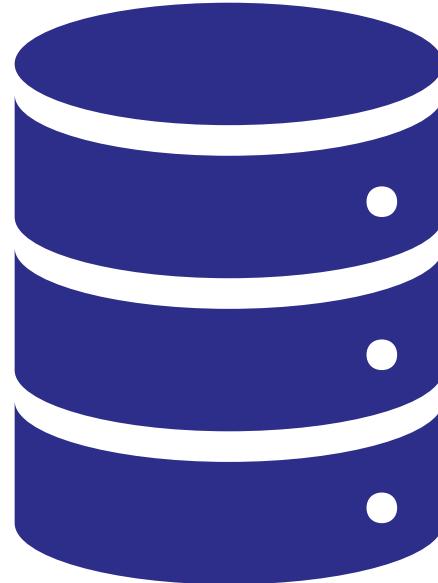
- Smart data placement enables better media utilization, lower WA
- Files can be mapped directly to erase units (zones)
 - **Minimal garbage collection**
 - **Lower write amplification**
- No longer a need to reserve over-provisioned media
 - **~7-28% more storage capacity**
- Enables QLC media for TLC use cases
- Reduces controller DRAM requirements



RocksDB

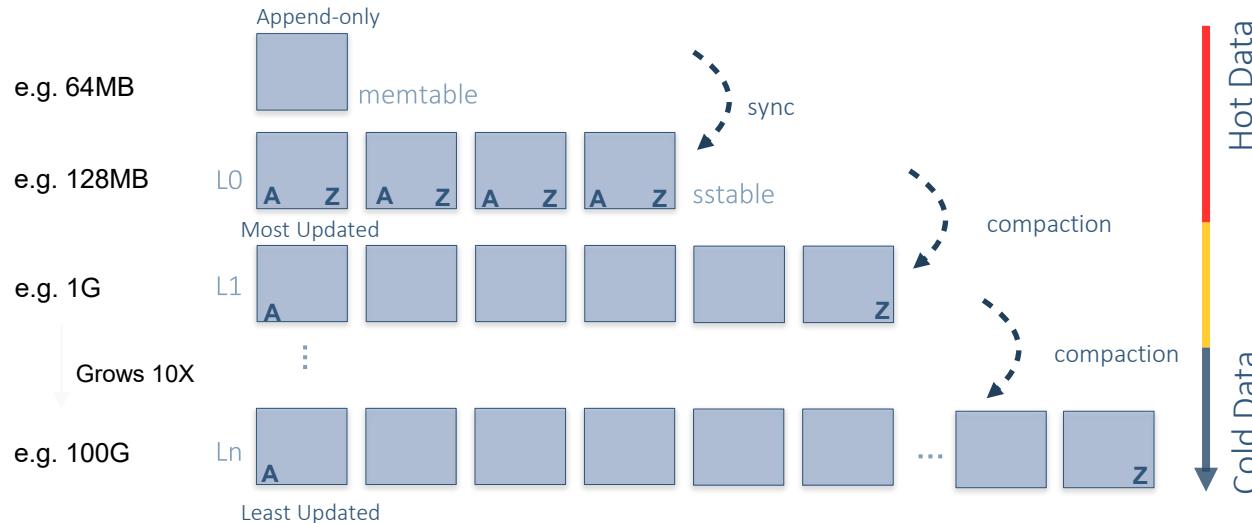
A great fit for ZNS adaptation!

- Widely used persistent key-value store for fast storage environments
- Log-structured, flash friendly
 - Clear separation of hot and cold data
- Open source
- Portable
- Active community
 - Up to date with latest storage stack improvements (e.g. io_uring)
- Pluggable storage backends



RocksDB on-disk data structures

Clear separation of hot and cold data



ZenFS

- Goal: end-to-end ZNS integration
 - Add native support for ZNS as a RocksDB FileSystem class
 - Use as much as possible of RocksDB data knowledge to do smart data placement
- Minimize write amplification
 - Avoid garbage collection
 - Minimize wear
 - Maximize write throughput
 - Improve read performance
- Minimize integration effort for users



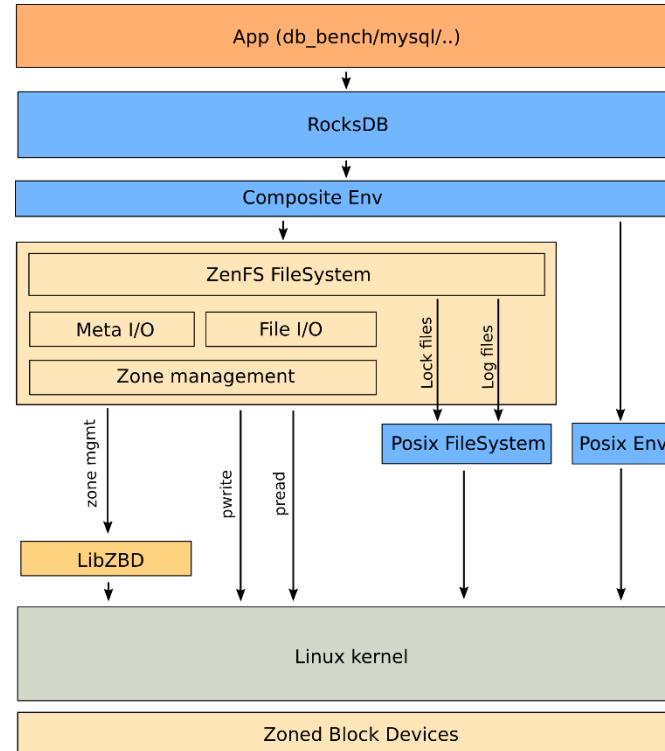
Architecture considerations

- Maximize reuse of existing components
- Leverage existing zone eco-system
- Keep complexity down to a minimum
- Flexibility in the right places
 - Disk IO path
 - Allocation algorithm



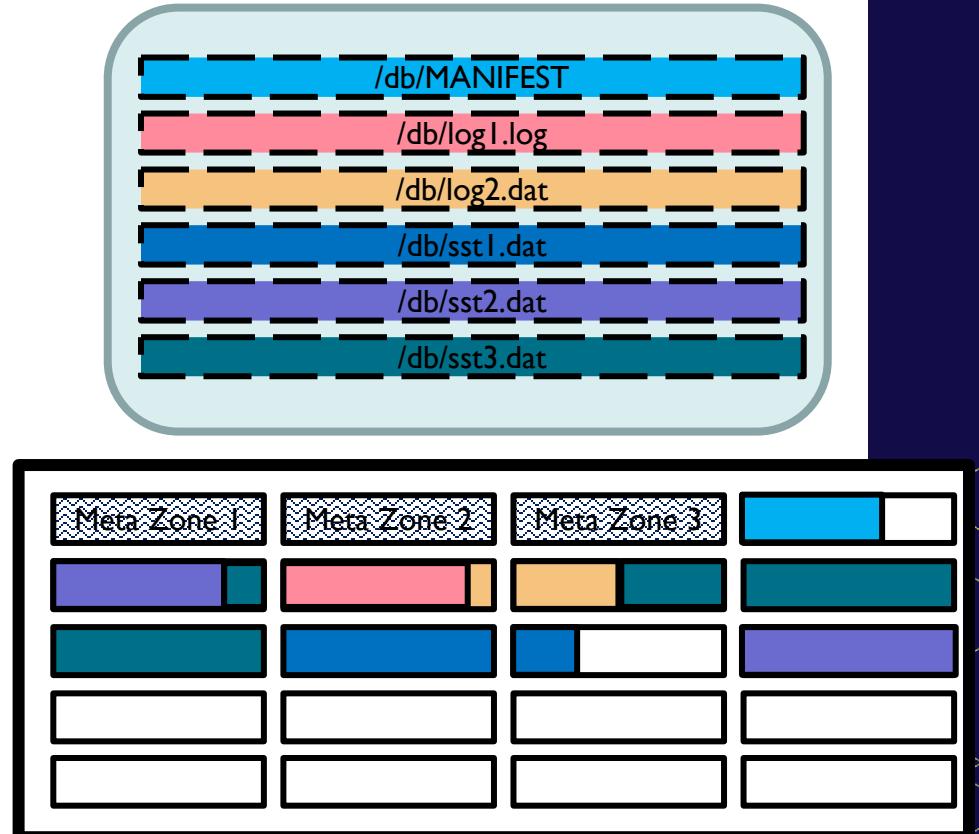
ZenFS Architecture

- Minimal impact on RocksDB itself
 - Adds a new ZenFS FileSystem class
 - **No changes** to the database code
- Stores all data and metadata files
 - Manifest
 - Write-ahead log
 - Sorted string tables (SSTs)
- Minimal external dependencies
 - Only relies on libzbd for zone operations
 - IO flows through conventional interfaces (pwrite/pread/libaio/io_uring..)
- ~ 2000 lines of code



ZenFS on-disk storage format

- File data is stored in a set of extents
 - Extents are a contiguous part of the address space **within a zone**
 - Files can be split over zones
 - Allows files of arbitrary size
 - Files can share zones
 - Enables full zone capacity usage
- ZenFs Metadata
 - File and global file system metadata
 - Stored incrementally in a rolling log
 - Occupies the first three zones



Zone Allocation

ZNS Constraints

- Sequential writes only
- Limited number of open zones
- Valid data must be evacuated before we can reuse a zone
 - Garbage collection is an option, but this leads to write amplification

Help from RocksDB

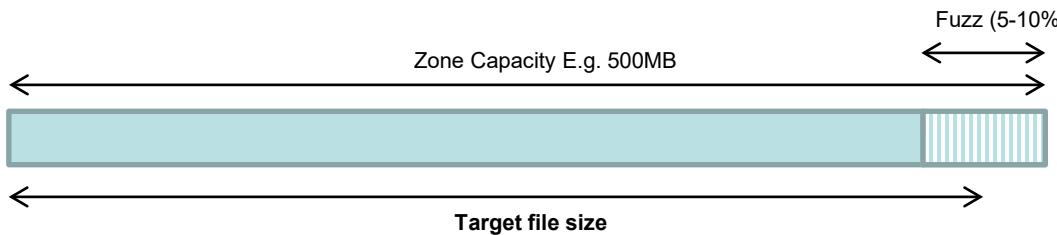
- Write life time hints
 - Metadata, WAL and SSTs
 - Different life time hints for SST levels
- File size hints
 - RocksDB provides file size hints before writing
- Append-only writes

ZenFS extent allocator

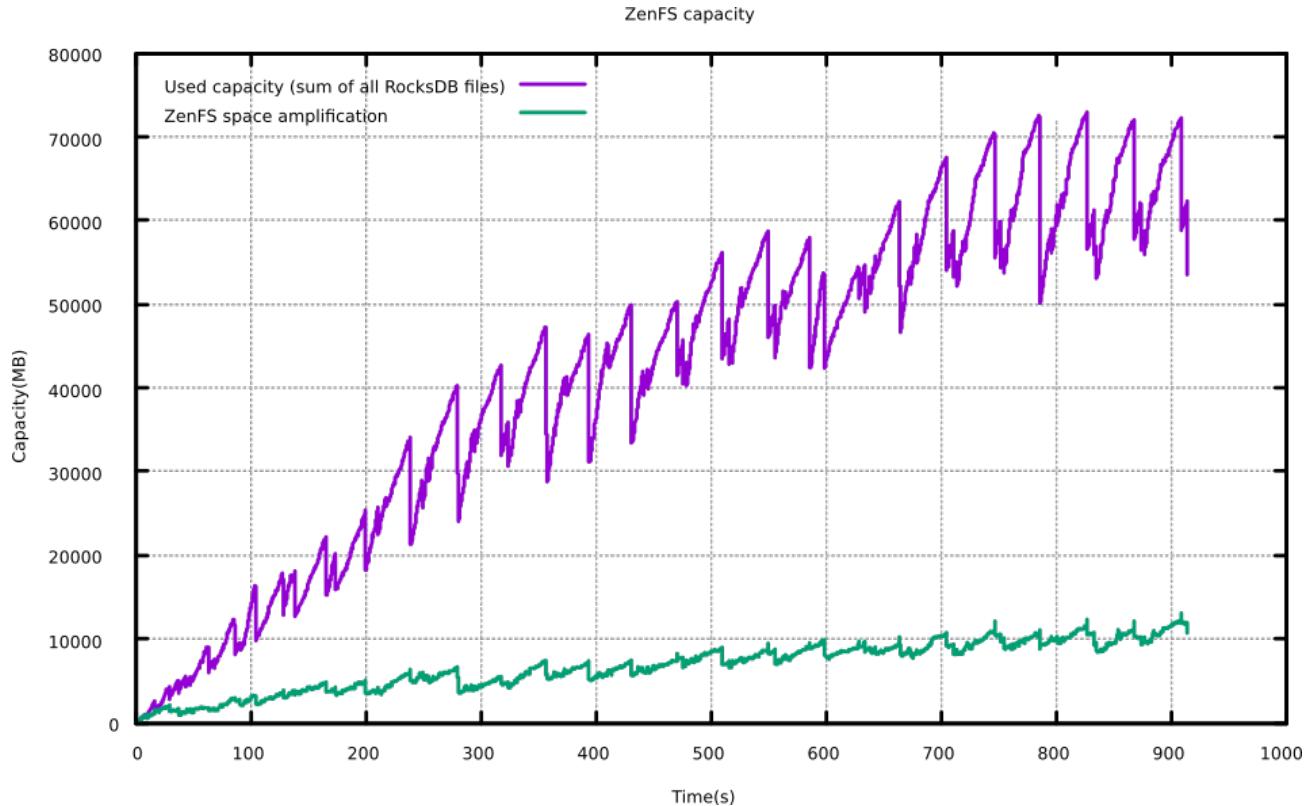
- Pick a zone for a new extent:
 - When a file is first written to
 - When there is no capacity left in the zone
- Current zone picking scheme:
 - If there is space left in a zone, fill it with warmer data to maximize capacity usage
 - Pick the first zone that fulfills: **Zone_lifetime > File_lifetime** OR is empty
 - If the picked zone is empty: **Zone_lifetime := File_lifetime**
- Files hold a write lock on the zone they are writing to
- Zones are marked as free when all files using the zone has been deleted
 - **No GC, No write amplification**

Fuzzy file size alignment

- The best capacity utilization is achieved if file sizes are aligned with zone capacity
- Target file size can be configured - but is just a hint to RocksDB
 - Allow for approximate/fuzzy zone capacity alignment
 - Set target file size in the middle of the fuzz area
 - Finish zones when they have reached the fuzz area

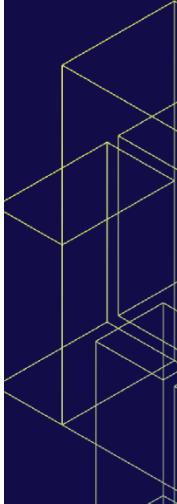


Space amplification



Write amplification = ~1.0

Allocation visualization



ZenFS Development

- Debugging/testing
 - db_bench, db_stress
 - gdb
- Tracing
 - A ZenFS log file is generated for each run on debug builds
 - /tmp/zenfs_<blockdevname>_* .log
 - Traces file operations and zone allocations
 - Great for debugging and evaluating extent allocation algorithms

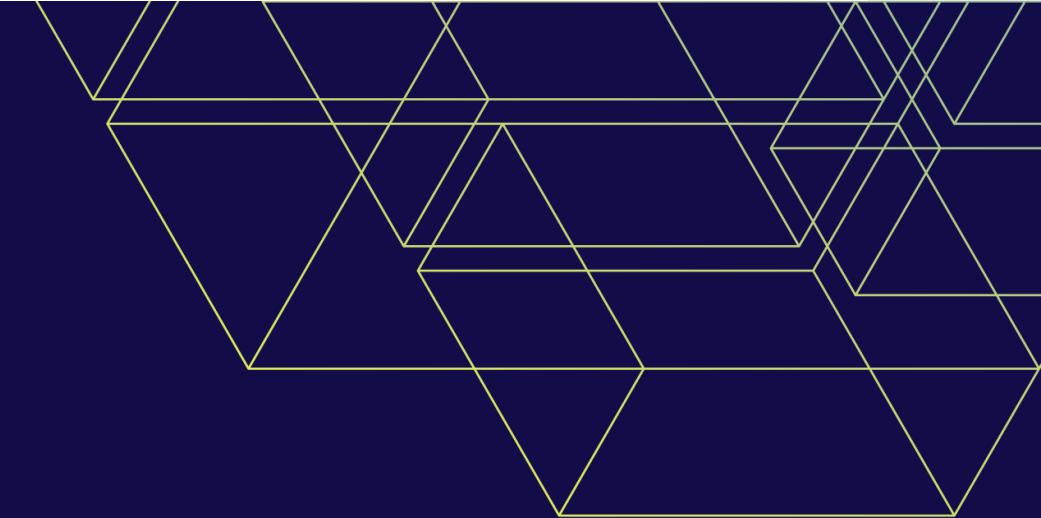


Current state

- Plumbing done
- Works with all Zoned Block Devices
 - ZNS SSDs, SMR Drives
 - Emulation/test devices: nullblk, QEMU
- Upstreaming in progress

Try it out!

- Easy to get started
 - Use a memory-backed nullblk as a zoned block device
- Latest code with instructions(see README.md):
 - <https://github.com/westerndigitalcorporation/rocksdb>
- For more information on zones and the eco-system
 - <https://zonedstorage.io>



Thanks!